

1 FPGA Implementation

1.1 Design Description

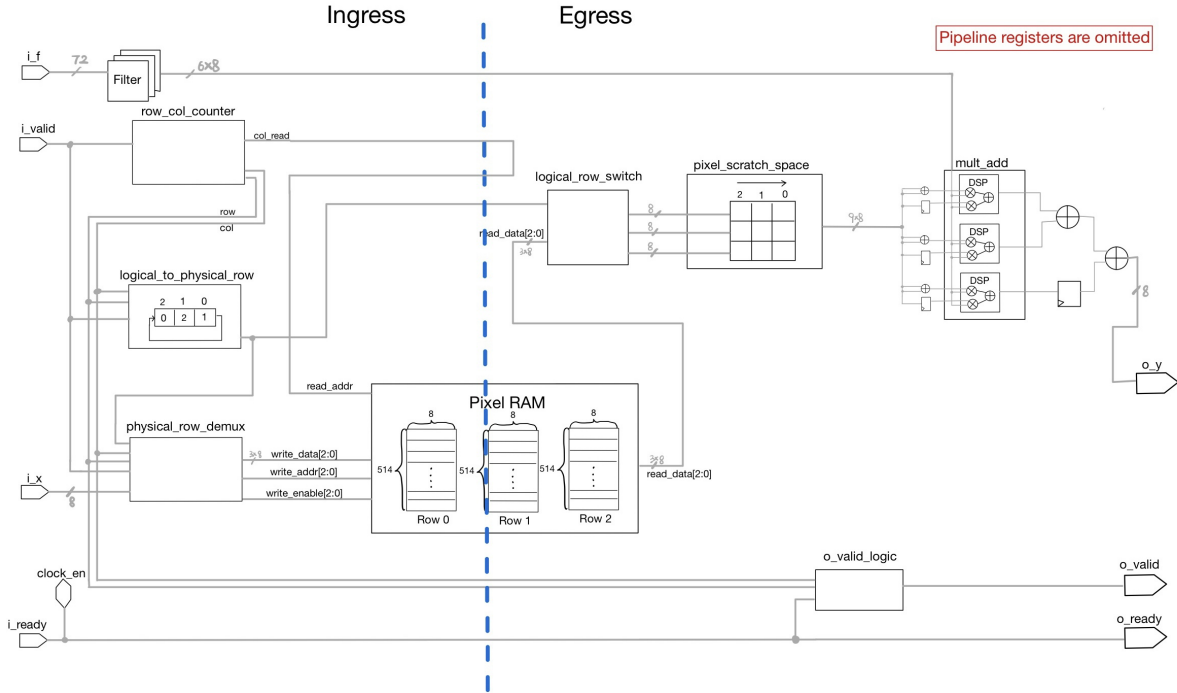


Figure 1: A simplified block diagram of the design, which runs at Restricted Fmax 549.45 MHz, Fmax 687.76 MHz, with a throughput of 1 output pixel per cycle

The design can be logically separated into two main stages, ingress and egress, with respect to the RAM that stores the input image pixels (pixel RAM).

Ingress path is responsible for keeping track of the column and row value of the input pixel, which is used as both the write address to the pixel RAM, and for determining the current convolution window (i.e., pointing to the bottom right corner of the convolution window). Another important responsibility for ingress path is to control the row shifting logic for the pixel RAM.

Pixel RAM is a 2D array of pixels, with a dimension of $3 \times (512 + 2) \times 8$. Only three rows are stored because that is the size of the convolution filter. Write is always done at the bottom row, and the row shifts upward when column counter reaches the bound. To save area, this shifting register file is implemented in BRAM, where the shifting is realized by shifting the mapping from logical row index to physical row index circularly, so that after the shifting, new incoming pixels are essentially overwriting the oldest row. Pixel RAM is implemented with three independent BRAM (M20K) blocks, each representing a row, with a shape of $(512 + 2) \times 8$. For ease of testing and maximal compatibility, the BRAM block

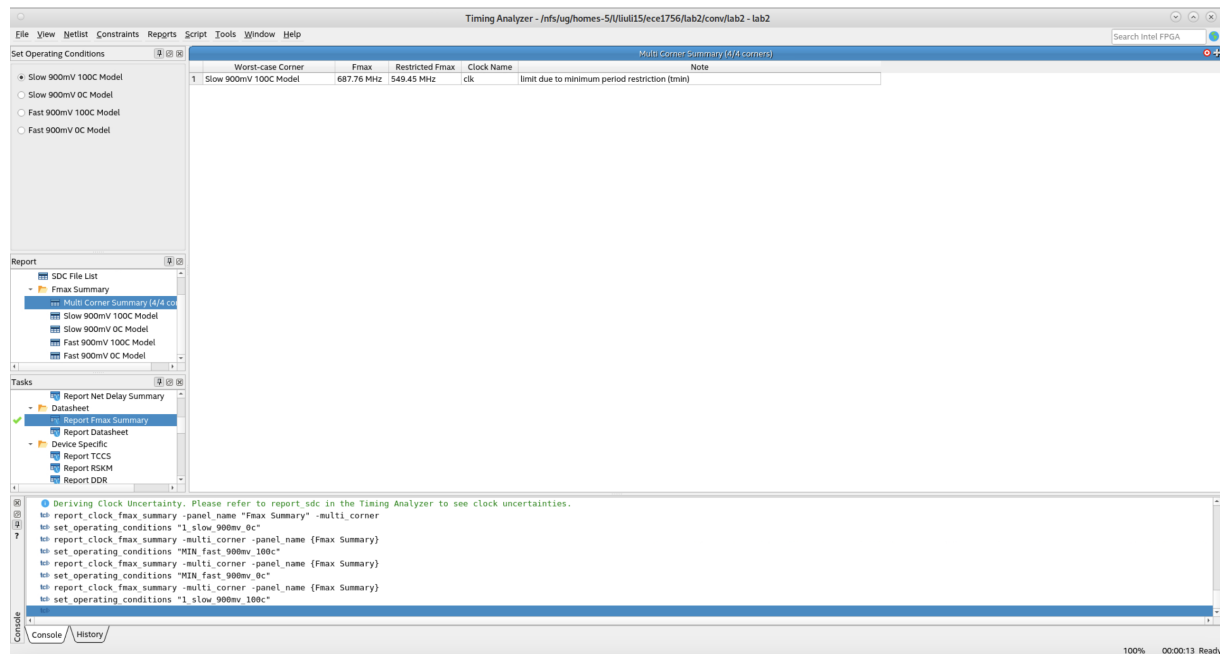


Figure 2: Fmax Summary

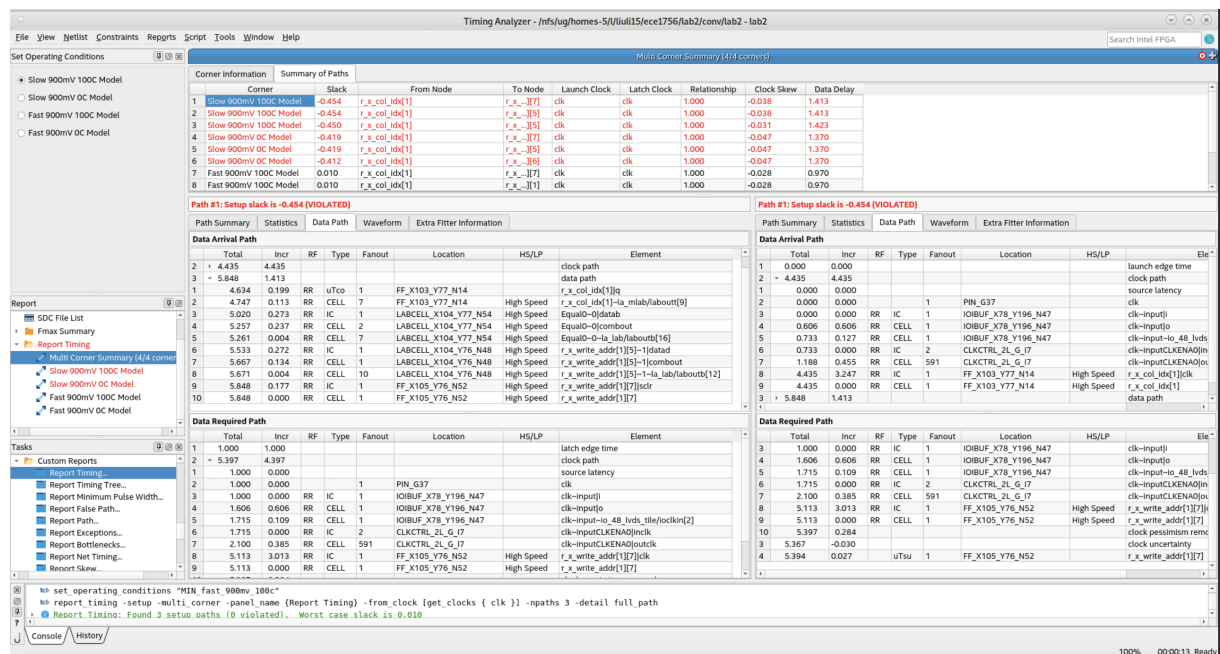


Figure 3: Critical Paths

Fitter Summary	
Search <<Filter>>	
Fitter Status	Successful - Mon Oct 31 22:29:35 2022
Quartus Prime Version	21.1.1 Build 850 06/23/2022 SJ Standard Edition
Revision Name	lab2
Top-level Entity Name	lab2
Family	Arria 10
Device	10AX115N1F45I1SG
Timing Models	Final
Logic utilization (in ALMs)	213 / 427,200 (< 1 %)
Total registers	564
Total pins	1 / 992 (< 1 %)
Total virtual pins	93
Total block memory bits	12,336 / 55,562,240 (< 1 %)
Total RAM Blocks	3 / 2,713 (< 1 %)
Total DSP Blocks	3 / 1,518 (< 1 %)
Total HSSI RX channels	0 / 48 (0 %)
Total HSSI TX channels	0 / 48 (0 %)
Total PLLs	0 / 112 (0 %)

Figure 4: Fitter Summary

is not instantiated from IP, but rather using an HDL template that can be inferred into BRAM (M20K) in simple dual-port mode by Quartus Synthesis, with read and write width of 8. With three independent BRAM (M20K) blocks, the pixel RAM is capable of doing a simultaneous read and write on each row in any clock cycle.

Egress path is responsible for performing the actual convolution computation for the convolution window that is pointed by the row and column counter aforementioned. For every window, it reads a single pixel for each row from the pixel RAM, and stores it into a scratch space that is 3×3 in size. The scratch space would shift horizontally to represent the fact that the convolution windows is also moving horizontally. So at any convolution window, only the right-most column is freshly read from the pixel RAM, while the other columns are cached from previous convolution windows. This effectively reduces the number of simultaneous read demand from 9 in a naive implementation, down to 3. To better exploit the fact that the filter is symmetric in x-direction, pixels at column 0 and column 2 are added before multiplying the filter, which effectively saves 3 multipliers in total. The total of 6 multipliers are divided into 3 sum of 2 18×18 multipliers groups, where each group can nicely fit into a DSP block on chip. Again, an HDL template that can be inferred to a fully registered DSP is used rather than instantiating from IP. Afterwards, a 2-level reduction (sum) tree is used to accumulate from the 3 DSP outputs. In total, 8 adders and 6 multipliers are used saving 3 multipliers from a naive implementation of 8 adders and 9 multipliers. A lot of

pipeline stages are inserted to this arithmetic chain to boost Fmax. The entire arithmetic logic runs continuously, but the output is valid only if the convolution window pointer is at valid location and has not been computed before.

1.2 FPGA Implementation Results

	Result
ALM Utilization	166 out of 427200
DSP Utilization	3 out of 1518
BRAM (M20K) Utilization	3 out of 2713
Maximum Operating Frequency (MHz)	Restricted Fmax 549.45 MHz, Fmax 687.76 MHz
Cycles for Test 7a (Hinton)	264222
Dynamic Power for one module @ maximum frequency (W)	0.048 W
Throughput of one module (GOPS)	9.267 GOPS
Throughput of full device (GOPS)	4689.198 GOPS
Total Power for full device (W)	10.188 W

Calculations

Dynamic Power for one module @ maximum frequency (W): $(DSP + M20K + LUT + FF + ClockNetwork) \div 50MHz \times RestrictedFmax = (0.08mW + 0.45mW + 0.20mW + 0.79mW + 2.82mW) \div 50MHz \times 549.45MHz$

Throughput of one module (GOPS): $512 \times 512 \times (9 + 8) \div (264222 \div 549.45MHz)$

Throughput of full device (GOPS): $506 \times 9.267GOPS$

Total Power for full device (W): $(DSP + M20K + LUT + FF) \div 50MHz \times RestrictedFmax \times 506 + ClockNetwork \div 50MHz \times RestrictedFmax + DeviceStatic(1704.56mW) + IO(0.19mW)$



1.3 Simulation Waveforms and Testbench Output

Refer to Figure 5 and Figure 6.

2 Efficiency Comparison

2.1 CPU Implementation

CPU convolution implementations are quite straight forward. Both single-threaded and multithreaded, and both basic and hand-vectorized versions loop for filter_id, output pixel



row and column respectively. Then a dot product is computed for the 3×3 convolution window involved. The basic implementation traverses through all 9 elements and accumulates the products, whereas the hand-vectorized version computes the dot product of the entire row in one shot (load, dot product and store) and then accumulate the the results of three rows together. Multi-threading is done using OpenMP directives, and is applied in the outermost filter_id dimension.

2.2 GPU Implementation

GPU convolution kernel executes in two stages. Firstly, each thread loads a few pixels of the current tile into shared memory. Then after all threads within the current block are done, each thread starts to calculate the dot product for the output pixel and filter_id assigned to it. The filters are stored in the GPU constant memory.

2.3 Runtime Results of CPU and GPU Implementations

No. of filters	Runtime (ms)			
	1	4	16	64
GPU	0.0151	0.0401	0.1374	0.5060
CPU (basic - no opt - 1 thread)	6.1022	24.306	97.235	389.02
CPU (vectorized - no opt - 1 thread)	3.1988	13.027	51.896	207.30
CPU (basic - O2 - 1 thread)	1.2414	4.9722	19.807	79.261
CPU (vectorized - O2 - 1 thread)	0.8441	3.3751	13.493	54.013
CPU (basic - O3 - 1 thread)	0.5334	2.1221	8.4699	33.843
CPU (vectorized - O3 - 1 thread)	0.8823	3.4746	13.898	55.689
CPU (basic - O3 - 4 threads)	0.5770	0.6108	1.2119	4.8529
CPU (vectorized - O3 - 4 threads)	0.9082	1.0459	1.8682	7.7334

When there is no compiler optimization, the vectorized version has a 2x performance than the basic version. When the compiler uses O2 optimization, the gap becomes smaller but still in favor of the vectorized version. Both versions have significant improvements comparing to the no opt version, but basic version gains more out of it suggesting that O2 optimizes both the dot product and iteration overheads for the basic version whereas the vectorized version only gains improvement from iteration overheads. This is even more obvious under O3 optimization, where the basic version has much better performance than the vectorized version and the performance of the vectorized version is the same as under O2 optimization. After turning on multi-threading, comparing to the single-threaded result (with O3 optimization), the relative performance gap between the basic version and vectorized version stays the same. It is worth pointing out the speed up comparing to single-threaded version is dependent on the number of filters, there is no speed up when there is only 1 filter, and

about 3.5x speed up for 4 filters. The speed up caps at around 8x for both 16 and 64 filters primarily due to the fact that the CPU under test has 8 cores.

2.4 Comparison of Results Between the 3 Compute Platforms

	Throughput (GOPS)	Power (W)	Energy Efficiency (GOPS/W)	Area Efficiency (GOPS/mm ²)
FPGA (20 nm)	4689.2 GOPS	10.188 W	460.27 GOPS/W	11.723 GOPS/mm ²
CPU (14 nm)	58.771 GOPS	65 W	0.904 GOPS/W	0.213 GOPS/mm ²
GPU (8 nm)	563.63 GOPS	220 W	2.562 GOPS/W	1.434 GOPS/mm ²
FPGA (scaled to 8nm)	7502.7 GOPS	6.875 W	1091.3 GOPS/W	75.027 GOPS/mm ²
CPU (scaled to 8nm)	73.464 GOPS	56.875 W	1.292 GOPS/W	0.532 GOPS/mm ²

Calculations

Scaled Throughput (GOPS): $OriginalThroughput \times ClockSpeedScaling$

Scaled Power: $OriginalPower \times ClockSpeedScaling \times Power@sameclockrateScaling$

Scaled Area: $OriginalArea \times AreaScaling$

When scaling to the same process technology, throughput wise, FPGA is one order ahead of GPU due primarily to that FPGA is programmed and dedicated to compute the convolution. One advantage with the FPGA is that it outputs a pixel at each clock cycle, where as GPU takes several. The GPU version is also a magnitude ahead of CPU, primarily due to the massive parallelism that convolution operation exhibits and that GPU can better exploit it.

Energy efficiency wise, FPGA is almost 1000x ahead of CPU and 400x ahead of GPU. This time, CPU and GPU are within the same magnitude due to them both being a general purpose computing platform, whereas the programmed FPGA is more application specific thus avoiding a lot of the overhead cost associated with generality. GPU is still ahead of CPU because of its wider SIMD, reducing overhead cost of instruction fetching and decoding.

Area efficiency wise, FPGA is almost 150x ahead of CPU and 50x ahead of GPU. The reason is again related to the overhead cost with generality, and overhead cost of instruction fetching and decoding between CPU and GPU.

3 Appendix

3.1 HDL

```

1  // This module implements 2D convolution between a 3x3 filter and a
   ↪ 512-pixel-wide image of any height.
2  // It is assumed that the input image is padded with zeros such that the
   ↪ input and output images have
3  // the same size. The filter coefficients are symmetric in the x-direction
   ↪ (i.e.  $f[0][0] = f[0][2]$ ,
4  //  $f[1][0] = f[1][2]$ ,  $f[2][0] = f[2][2]$  for any filter  $f$ ) and their values
   ↪ are limited to integers
5  // (but can still be positive or negative). The input image is grayscale
   ↪ with 8-bit pixel values ranging
6  // from 0 (black) to 255 (white).
7  module lab2 (
8      input  clk,                // Operating clock
9      input  reset,              // Active-high reset signal
   ↪ (reset when set to 1)
10     input  [71:0] i_f,          // Nine 8-bit signed convolution
   ↪ filter coefficients in row-major format (i.e.  $i_f[7:0]$  is  $f[0][0]$ ,
   ↪  $i_f[15:8]$  is  $f[0][1]$ , etc.)
11     input  i_valid,             // Set to 1 if input pixel
   ↪ is valid
12     input  i_ready,            // Set to 1 if consumer
   ↪ block is ready to receive a new pixel
13     input  [7:0] i_x,          // Input pixel value (8-bit
   ↪ unsigned value between 0 and 255)
14     output o_valid,            // Set to 1 if output pixel
   ↪ is valid
15     output o_ready,            // Set to 1 if this block is
   ↪ ready to receive a new pixel
16     output [7:0] o_y           // Output pixel value (8-bit
   ↪ unsigned value between 0 and 255)
17 );
18
19 localparam FILTER_SIZE = 3;    // Convolution filter dimension (i.e.
   ↪ 3x3)
20 localparam PIXEL_DATAW = 8;    // Bit width of image pixels and filter
   ↪ coefficients (i.e. 8 bits)
21
22 integer col, row, i; // variables to use in the for loop

```



```
23
24 // The following code is intended to show you an example of how to use
    ↳ paramaters and
25 // for loops in SytemVerilog. It also arrages the input filter
    ↳ coefficients for you
26 // into a nicely-arranged and easy-to-use 2D array of registers. However,
    ↳ you can ignore
27 // this code and not use it if you wish to.
28
29 logic signed [PIXEL_DATAW-1:0] r_f [FILTER_SIZE-1:0][FILTER_SIZE-1:0]; // 2D
    ↳ array of registers for filter coefficients
30 always_ff @ (posedge clk) begin
31     // If reset signal is high, set all the filter coefficient
    ↳ registers to zeros
32     // We're using a synchronous reset, which is recommended style for
    ↳ recent FPGA architectures
33     if(reset)begin
34         for(row = 0; row < FILTER_SIZE; row = row + 1) begin
35             for(col = 0; col < FILTER_SIZE; col = col + 1) begin
36                 r_f[row][col] <= 0;
37             end
38         end
39         // Otherwise, register the input filter coefficients into the 2D
    ↳ array signal
40         end else begin
41             for(row = 0; row < FILTER_SIZE; row = row + 1) begin
42                 for(col = 0; col < FILTER_SIZE; col = col + 1) begin
43                     // Rearrange the 72-bit input into a 3x3
    ↳ array of 8-bit filter coefficients.
44                     // signal[a +: b] is equivalent to
    ↳ signal[a+b-1 : a]. You can try to plug in
45                     // values for col and row from 0 to 2, to
    ↳ understand how it operates.
46                     // For example at row=0 and col=0:
    ↳ r_f[0][0] = i_f[0+:8] = i_f[7:0]
47                     // at row=0 and col=1:
    ↳ r_f[0][1] = i_f[8+:8] = i_f[15:8]
48                     r_f[row][col] <= i_f[(row * FILTER_SIZE *
    ↳ PIXEL_DATAW)+(col * PIXEL_DATAW) +: PIXEL_DATAW];
49                 end
50             end
51         end
```

```
52 end
53
54 // Start of your code
55 logic enable;
56 assign enable = i_ready;
57
58 // *****
59 // INGRESS
60 // *****
61 // If pipelined, need to pipeline i_valid
62
63 // Logics for buffer of x
64 localparam IMAGE_WIDTH = 512;
65 localparam R_X_ROWS = FILTER_SIZE; // Always store 3 rows of i_x
66 localparam R_X_COL_WIDTH = IMAGE_WIDTH + 2;
67 localparam R_X_COL_ADDR_WIDTH = 10;
68
69
70 // Pixel RAM
71 // RAM input, need to be registered except for r_x_read_addr
72 logic [R_X_COL_ADDR_WIDTH-1:0] r_x_write_addr [R_X_ROWS-1:0]; // 0..511+2
73   ↪ (10 bit)
74 logic r_x_write_enable [R_X_ROWS-1:0];
75 logic unsigned [PIXEL_DATAW-1:0] r_x_write_data [R_X_ROWS-1:0];
76 // Registered inside module, read a pixel from 3 rows
77 logic [R_X_COL_ADDR_WIDTH-1:0] r_x_read_addr;
78 // RAM output
79 logic unsigned [PIXEL_DATAW-1:0] r_x_read_data [R_X_ROWS-1:0]; // is
80   ↪ registered inside module, 3 rows of 1 pixel
81 pixelram pixel_ram
82 (
83     .clk(clk),
84     .reset(reset),
85     .enable(enable),
86     // RAM input, unregistered inside module
87     .i_write_addr(r_x_write_addr),
88     .i_write_enable(r_x_write_enable),
89     .i_write_data(r_x_write_data),
90     .i_read_addr(r_x_read_addr),
91     // RAM output
92     .o_read_data(r_x_read_data)
93 );
```

```

92
93
94 // Registers
95 // 0: [] [] [] [] [] [] [] [] 512 + 2
96 // 1: [] [] [] [] [] [] [] [] 512 + 2
97 // 2: [] [] [] [] [] [] [] [] 512 + 2
98 logic unsigned [1:0] r_x_row_logical_idx; // Count from 0 to R_X_ROWS - 1
    ↪ (incl), logical order, not necessarily physical
99 logic unsigned [R_X_COL_ADDR_WIDTH-1:0] r_x_col_idx; // Count from 0 to
    ↪ R_X_COL_WIDTH (incl)
100 logic unsigned [R_X_COL_ADDR_WIDTH-1:0] r_x_col_idx_read_addr_adjusted; //
    ↪ Count from 0 to R_X_COL_WIDTH (incl)
101 // Count from 0 to R_X_ROWS - 1 (incl), physical order
102 logic unsigned [R_X_ROWS-1:0][1:0] r_x_row_logical_to_physical_index;
103
104 // INGRESS: Stage -1
105 always_ff @ (posedge clk) begin
106     if(reset) begin
107         r_x_row_logical_idx <= 0;
108         r_x_col_idx <= 0;
109         r_x_col_idx_read_addr_adjusted <= 0;
110
111         for(i = 0; i < R_X_ROWS; i = i + 1) begin
112             r_x_row_logical_to_physical_index[i] <= i;
113         end
114
115         for(row = 0; row < R_X_ROWS; row = row + 1) begin
116             r_x_write_addr[row] <= 0;
117             r_x_write_data[row] <= 0;
118             r_x_write_enable[row] <= 0;
119         end
120     end else if (enable) begin
121         // Do not write by default
122         for(row = 0; row < R_X_ROWS; row = row + 1) begin
123             r_x_write_addr[row] <= 0;
124             r_x_write_data[row] <= 0;
125             r_x_write_enable[row] <= 0;
126         end
127         r_x_col_idx_read_addr_adjusted <= 0;
128         if(i_valid) begin
129             if(r_x_col_idx == R_X_COL_WIDTH) begin

```

```
130          // Load input pixel to a new row at the
↪  current logical idx 0 (R_X_COL_WIDTH implies 0),
131          // which would be discarded, then
↪  reused/overwritten as the new logical idx 2 in the next cycle
132          r_x_write_addr[r_x_row_logical_to_physical_index[0]]
↪  <= 0;
133          r_x_write_data[r_x_row_logical_to_physical_index[0]]
↪  <= i_x;
134          r_x_write_enable[r_x_row_logical_to_physical_index[0]]
↪  <= 1;
135
136          // Do the row shifting logic at the first
↪  input of the new row,
137          // rather than at the last input of the old
↪  row (will have conflict
138          // in writing old row and shifting old row)
139
140          // Instead of shifting the actual data,
↪  shift the mapping from logical index to physical index
141          // Shift the mapping, upward
↪  (idx[0]->idx[2])
142          r_x_row_logical_to_physical_index <=
↪  {r_x_row_logical_to_physical_index[0],
↪  r_x_row_logical_to_physical_index[R_X_ROWS-1:1]};
143
144          // Reset r_x_col_idx if necessary,
↪  continuing at idx 1.
145          // Skipping idx 0 because we are at idx 0
↪  currently
146          r_x_col_idx <= 1;
147          r_x_col_idx_read_addr_adjusted <= 0;
148
149          // Increment r_x_row_logical_idx_ipipelined
↪  only when r_x_row_logical_idx_ipipelined is 0 or 1,
150          // so that r_x_row_logical_idx_ipipelined
↪  will reach to 2 in steady state
151          if(r_x_row_logical_idx < R_X_ROWS - 1) begin
152              r_x_row_logical_idx <=
↪  r_x_row_logical_idx + 1;
153          end
154          end else begin
155              // Load data at logical idx 2
```

```
156         r_x_write_addr[r_x_row_logical_to_physical_index[R_X_ROW
↪      <= r_x_col_idx;
157         r_x_write_data[r_x_row_logical_to_physical_index[R_X_ROW
↪      <= i_x;
158         r_x_write_enable[r_x_row_logical_to_physical_index[R_X_R
↪      <= 1;
159
160         // Increment r_x_col_idx
161         r_x_col_idx <= r_x_col_idx + 1;
162
163         // Increment r_x_col_idx_read_addr_adjusted
164         r_x_col_idx_read_addr_adjusted <=
↪      r_x_col_idx;
165         end
166     end
167 end
168
169 // Pipeline registers for ingress
170 logic unsigned [1:0] r_x_row_logical_idx_ipipelined;
171 logic unsigned [R_X_COL_ADDR_WIDTH-1:0] r_x_col_idx_ipipelined;
172 logic unsigned [R_X_ROWS-1:0][1:0]
↪      r_x_row_logical_to_physical_index_ipipelined;
173 // INGRESS: Stage 0, to match registered RAM write
174 always_ff @ (posedge clk) begin
175     if (enable) begin
176         r_x_row_logical_idx_ipipelined <= r_x_row_logical_idx;
177         r_x_col_idx_ipipelined <= r_x_col_idx;
178         r_x_row_logical_to_physical_index_ipipelined <=
↪      r_x_row_logical_to_physical_index;
179     end
180 end
181
182 // *****
183 // EGRESS
184 // *****
185
186 // Pipeline registers for egress
187 localparam NUM_EGRESS_STAGE = 9;
188 logic unsigned [NUM_EGRESS_STAGE-1:0] [R_X_COL_ADDR_WIDTH-1:0]
↪      r_x_col_idx_epipelined;
189 logic unsigned [NUM_EGRESS_STAGE-1:0] [1:0] r_x_row_logical_idx_epipelined;
```

```
191 logic unsigned [1:0] [R_X_ROWS-1:0][1:0]
    ↳ r_x_row_logical_to_physical_index_epipelined; // Not needed for full
    ↳ pipeline stage
192 always_ff @ (posedge clk) begin
193     if(enable) begin
194         r_x_col_idx_epipelined <=
    ↳ {r_x_col_idx_epipelined[NUM_EGRESS_STAGE-2:0], r_x_col_idx_ipipelined};
195         r_x_row_logical_idx_epipelined <=
    ↳ {r_x_row_logical_idx_epipelined[NUM_EGRESS_STAGE-2:0],
    ↳ r_x_row_logical_idx_ipipelined};
196         r_x_row_logical_to_physical_index_epipelined <=
    ↳ {r_x_row_logical_to_physical_index_epipelined[0:0],
    ↳ r_x_row_logical_to_physical_index_ipipelined};
197     end
198 end
199
200 // Logics for convolution core
201 // Computation
202
203 // EGRESS: Stage -1
204 always_comb begin
205     r_x_read_addr = r_x_col_idx_read_addr_adjusted;
206 end
207
208 // EGRESS: Stage 0, 1
209 // Signed x unsigned gets unsigned, which is not what we intend.
210 // So convert unsigned to signed by treating unsigned number as positive
    ↳ (by adding a 0 to msb)
211 logic signed [FILTER_SIZE-1:0] [PIXEL_DATAW:0] r_mult_i_pixel
    ↳ [R_X_ROWS-1:0];
212 logic [PIXEL_DATAW:0] r_x_read_data_reg [R_X_ROWS-1:0];
213 always_ff @ (posedge clk) begin
214     if(enable) begin
215         for(row=0; row<R_X_ROWS; row=row+1) begin
216             r_x_read_data_reg[row] <= r_x_read_data[row];
217             r_mult_i_pixel[row] <= {{1'b0,
    ↳ r_x_read_data_reg[r_x_row_logical_to_physical_index_epipelined[1][row]]},
    ↳ r_mult_i_pixel[row][FILTER_SIZE-1:1]}};
218         end
219     end
220 end
221
```

```
222 // Multiplication
223 // EGRESS: Stage 2, 3, 4, 5
224 logic signed [FILTER_SIZE-1:0] [2*PIXEL_DATAW-1:0] sums_stage_0;
225 genvar gen_row;
226 generate
227     for(gen_row=0; gen_row<FILTER_SIZE; gen_row=gen_row+1) begin: mult
228         mult8x8p8x8 m0(
229             .clk(clk),
230             .enable(enable),
231             .i_filtera(r_f[gen_row][0]),
232             .i_pixela0(r_mult_i_pixel[gen_row][0]),
233             .i_pixela1(r_mult_i_pixel[gen_row][2]),
234             .i_filterb(r_f[gen_row][1]),
235             .i_pixelb(r_mult_i_pixel[gen_row][1]),
236             .o_res(sums_stage_0[gen_row])
237         );
238     end
239 endgenerate
240
241 // Reduction tree
242 // EGRESS: Stage 6
243 logic signed [FILTER_SIZE-1:0] [2*PIXEL_DATAW-1:0] sums_stage_0_reg;
244 always_ff @ (posedge clk) begin
245     if(enable) begin
246         sums_stage_0_reg <= sums_stage_0;
247     end
248 end
249
250 logic signed [2*PIXEL_DATAW-1:0] sums_stage_1 [1:0];
251 always_comb begin
252     sums_stage_1[0] = sums_stage_0_reg[0] + sums_stage_0_reg[1];
253     sums_stage_1[1] = sums_stage_0_reg[2];
254 end
255
256 logic signed [2*PIXEL_DATAW-1:0] sums_stage_2;
257 always_comb begin
258     sums_stage_2 = sums_stage_1[0] + sums_stage_1[1];
259 end
260 // EGRESS: Stage 7
261 logic signed [2*PIXEL_DATAW-1:0] sums_stage_2_reg;
262 always_ff @ (posedge clk) begin
263     if(enable)begin
```

```
264         sums_stage_2_reg <= sums_stage_2;
265     end
266 end
267
268 logic unsigned [PIXEL_DATAW-1:0] y;
269 always_comb begin
270     if(sums_stage_2_reg>255) begin
271         y = 255;
272     end else if (sums_stage_2_reg<0) begin
273         y = 0;
274     end else begin
275         y = sums_stage_2_reg[PIXEL_DATAW-1:0];
276     end
277 end
278
279 // Output interface logics
280 // EGRESS: Stage 8
281 logic unsigned [PIXEL_DATAW-1:0] r_y;
282 logic r_y_valid;
283 logic unsigned [R_X_COL_ADDR_WIDTH-1:0] r_x_col_idx_prev;
284 always_ff @ (posedge clk) begin
285     if(reset) begin
286         r_y <= 0;
287         r_y_valid <= 0;
288         r_x_col_idx_prev <= 0;
289     end else if(enable) begin
290         r_x_col_idx_prev <=
↪ r_x_col_idx_epipelined[NUM_EGRESS_STAGE-1];
291         // By the time r_x_col_idx is 3, pixel at idx 2 is already
↪ written with i_x
292         if(r_x_col_idx_prev !=
↪ r_x_col_idx_epipelined[NUM_EGRESS_STAGE-1] &&
293         r_x_col_idx_epipelined[NUM_EGRESS_STAGE-1] >=
↪ FILTER_SIZE &&
294         r_x_row_logical_idx_epipelined[NUM_EGRESS_STAGE-1]
↪ == R_X_ROWS - 1) begin
295             r_y <= y;
296             r_y_valid <= 1;
297         end else begin
298             r_y <= 0;
299             r_y_valid <= 0;
300         end
end
```



```
301         end
302     end
303
304     assign o_y = r_y;
305     // Ready for inputs as long as receiver is ready for outputs
306     assign o_ready = i_ready;
307     assign o_valid = r_y_valid & i_ready;
308
309     // End of your code
310
311 endmodule
312
313 /*****
314
315 // Multiplier module for 8x8 multiplications + 8x8 multiplications
316 module mult8x8p8x8 (
317     input clk,
318     input enable,
319     input signed [7:0] i_filtera,
320     input signed [8:0] i_pixela0,
321     input signed [8:0] i_pixela1,
322     input signed [7:0] i_filterb,
323     input signed [8:0] i_pixelb,
324     output logic signed [15:0] o_res
325 );
326
327 // Pipeline 0
328 logic signed [8:0] i_pixela0_reg, i_pixela1_reg, i_pixelb_reg;
329 logic signed [7:0] i_filtera_reg, i_filterb_reg;
330 always_ff @(posedge clk) begin
331     if(enable)begin
332         i_pixela0_reg <= i_pixela0;
333         i_pixela1_reg <= i_pixela1;
334         i_pixelb_reg <= i_pixelb;
335         i_filtera_reg <= i_filtera;
336         i_filterb_reg <= i_filterb;
337     end
338 end
339
340 // Pipeline 1
341 logic signed [9:0] i_pixela01_reg_reg;
342 logic signed [8:0] i_pixelb_reg_reg;
```

```
343 logic signed [7:0] i_filtera_reg_reg, i_filterb_reg_reg;
344 always_ff @ (posedge clk) begin
345     if(enable) begin
346         i_pixela01_reg_reg <= i_pixela0_reg + i_pixela1_reg;
347         i_pixelb_reg_reg <= i_pixelb_reg;
348         i_filtera_reg_reg <= i_filtera_reg;
349         i_filterb_reg_reg <= i_filterb_reg;
350     end
351 end
352
353 // Pipeline 2
354 logic signed [9:0] i_pixela01_reg_reg_reg;
355 logic signed [8:0] i_pixelb_reg_reg_reg;
356 logic signed [7:0] i_filtera_reg_reg_reg, i_filterb_reg_reg_reg;
357 always_ff @ (posedge clk) begin
358     if(enable) begin
359         i_pixela01_reg_reg_reg <= i_pixela01_reg_reg;
360         i_pixelb_reg_reg_reg <= i_pixelb_reg_reg;
361         i_filtera_reg_reg_reg <= i_filtera_reg_reg;
362         i_filterb_reg_reg_reg <= i_filterb_reg_reg;
363     end
364 end
365
366 // Pipeline 3
367 always_ff @ (posedge clk) begin
368     if(enable) begin
369         o_res <= i_pixela01_reg_reg_reg * i_filtera_reg_reg_reg +
↪ i_pixelb_reg_reg_reg * i_filterb_reg_reg_reg;
370     end
371 end
372 endmodule
373
374 /*****
375
376 module pixelram #
377 (
378     parameter FILTER_SIZE = 3,
379     parameter PIXEL_DATAW = 8,
380     parameter IMAGE_WIDTH = 512,
381     parameter R_X_ROWS = FILTER_SIZE,
382     parameter R_X_COL_ADDR_WIDTH = 10
383 )
```

```

384 (
385     input clk,
386     input reset,
387     input enable,
388     // RAM input, unregistered inside the module except for
    ↪ i_read_addr
389     input [R_X_COL_ADDR_WIDTH-1:0] i_write_addr [R_X_ROWS-1:0],
390     input i_write_enable [R_X_ROWS-1:0],
391     input unsigned [PIXEL_DATAW-1:0] i_write_data [R_X_ROWS-1:0],
392     input [R_X_COL_ADDR_WIDTH-1:0] i_read_addr, // registered inside
    ↪ the module
393     // RAM output
394     output unsigned [PIXEL_DATAW-1:0] o_read_data [R_X_ROWS-1:0] //
    ↪ registered, 3 rows of 1 pixel
395 );
396     // Wrap as RAM
397     // 0: [] [] [] [] [] [] [] []          512 + 2
398     // 1: [] [] [] [] [] [] [] []          512 + 2
399     // 2: [] [] [] [] [] [] [] []          512 + 2
400     genvar gen_row;
401     generate
402         for(gen_row=0; gen_row<R_X_ROWS; gen_row=gen_row+1) begin:
    ↪ pixel_ram_row
403             pixelrowram pixel_row_ram
404             (
405                 .clk(clk),
406                 .reset(reset),
407                 .enable(enable),
408                 // RAM input, unregistered
409                 .i_write_addr(i_write_addr[gen_row]),
410                 .i_write_enable(i_write_enable[gen_row]),
411                 .i_write_data(i_write_data[gen_row]),
412                 .i_read_addr(i_read_addr), // read 1 pixel
413                 // RAM output
414                 .o_read_data(o_read_data[gen_row]) //
    ↪ registered, 1 pixel
415             );
416         end
417     endgenerate
418 endmodule
419
420 /*****

```

```
421
422 module pixelrowram #
423 (
424     parameter PIXEL_DATAW = 8,
425     parameter IMAGE_WIDTH = 512,
426     parameter R_X_COL_WIDTH = IMAGE_WIDTH + 2,
427     parameter R_X_COL_ADDR_WIDTH = 10
428 )
429 (
430     input clk,
431     input reset,
432     input enable,
433     // RAM input, unregistered inside the module except for
    ↪ i_read_addr
434     input [R_X_COL_ADDR_WIDTH-1:0] i_write_addr,
435     input i_write_enable,
436     input unsigned [PIXEL_DATAW-1:0] i_write_data,
437     input [R_X_COL_ADDR_WIDTH-1:0] i_read_addr, // registered inside
    ↪ the module
438     // RAM output
439     output logic unsigned [PIXEL_DATAW-1:0] o_read_data // registered
440 );
441     // Wrap as RAM
442     // 0: [] [] [] [] [] [] [] [] 512 + 2
443     // 2D array of registers for input pixels, row major
444     // set_global_assignment -name
    ↪ ADD_PASS_THROUGH_LOGIC_TO_INFERRED_RAMs OFF
445     logic unsigned [PIXEL_DATAW-1:0] mem [R_X_COL_WIDTH-1:0];
446
447     logic [R_X_COL_ADDR_WIDTH-1:0] i_read_addr_reg;
448
449     integer i;
450     initial begin
451         for (i=0; i<R_X_COL_WIDTH; i=i+1) begin
452             mem[i] = 0;
453         end
454     end
455
456     always_ff @ (posedge clk) begin
457         if(enable) begin
458             i_read_addr_reg <= i_read_addr;
459         end
    
```

```
460         end
461
462         always_ff @ (posedge clk) begin
463             if(enable) begin
464                 if(i_write_enable) begin
465                     mem[i_write_addr] <= i_write_data;
466                 end
467                 o_read_data <= mem[i_read_addr_reg];
468             end
469         end
470     endmodule
```